

Fastai Lesson 11a review

Joseph Catanzarite
Fastai Deep Learning From The Foundations
TWiML Study Group Meetup
Saturday, 10/12/2019

Overview

Today:

Start Lesson 11

Notebooks

- 07a_lsuv
- 08_data_block

Next meeting:

Continue in Lesson 11

Listen to / watch at least the first half of the Lesson 11 video, through the DataBlocks API part.

Notebooks

- Finish 08_data_block
- Start 09_optimizers

Fastai Forum post:

Questions about RunningBatchNorm
class in Lesson 10

- <https://forums.fast.ai/t/questions-about-runningbatchnorm-in-lesson-10/56331>



The screenshot shows a Jupyter Notebook interface with a browser window at the top displaying the URL `localhost:8889/notebooks/dev_course/dl2/07_batchnorm.ipynb#`. The notebook has a menu bar with options: File, Edit, View, Insert, Cell, Kernel, Widgets, and Help. A 'Not Trusted' warning is visible in the top right corner of the notebook area. The main content of the notebook is a Python class definition for `RunningBatchNorm`.

Simplified RunningBatchNorm

It turns out we don't actually need to debias - because, for instance, dividing a debiased sum by a debiased count is the same as dividing a *biased* sum by a *biased* count! So we can remove all the debiasing stuff and end up with a simpler class. Also, we should save `eps` as a buffer since it impacts the calculation. (Thanks to Stas Bekman for noticing these.) Also we can slightly change the final calculation in `forward` with one that uses `factor` and `offset` to reduce the amount of broadcasting required. (Thanks to Tom Viehmann for this suggestion.)

```
In [ ]: #export
class RunningBatchNorm(nn.Module):
    def __init__(self, nf, mom=0.1, eps=1e-5):
        super().__init__()
        self.mom, self.eps = mom, eps
        self.mults = nn.Parameter(torch.ones(nf,1,1))
        self.adds = nn.Parameter(torch.zeros(nf,1,1))
        self.register_buffer('sums', torch.zeros(1,nf,1,1))
        self.register_buffer('sqrs', torch.zeros(1,nf,1,1))
        self.register_buffer('count', tensor(0.))
        self.register_buffer('factor', tensor(0.))
        self.register_buffer('offset', tensor(0.))
        self.batch = 0

    def update_stats(self, x):
        bs,nc,*_ = x.shape
        self.sums.detach_()
```

MORE VIDEOS



2:18 / 2:19:11



Fastai Forum post re: Simplified RunningBatchNorm in Lesson 11

- <https://forums.fast.ai/t/simplified-runningbatchnorm-in-lesson-11/56340>

Fastai Forum post:

Question on the Implementation of LSUV

- In the paper “All you need is a good init” the authors say that their Algorithm 1 is intended to follow a certain orthonormal initialization procedure: “We thus extend the orthonormal initialization Saxe et al. (2014) to an iterative procedure, described in Algorithm 1. Saxe et al. (2014) could be implemented in two steps. First, fill the weights with Gaussian noise with unit variance. Second, decompose them to orthonormal basis with QR or SVD decomposition and replace weights with one of the components.”
- Question @Sylvain, @stas or anyone else who knows: why did Fastai omit the Saxe et al. (2014) orthonormal initialization procedure?
- <https://forums.fast.ai/t/lesson-11-discussion-and-wiki/43406/366>

Note on a minor tweak that 'fixes' the LSUV normalization algorithm

- <https://forums.fast.ai/t/lesson-11-discussion-and-wiki/43406/365>

Running the Lesson 11 notebook 08_data_block.ipynb

- <https://forums.fast.ai/t/running-the-lesson-11-notebook-08-data-block-ipynb/56346>